

Illuminating Cognizance: A Comprehensive Look Into Major Power Outages in the U.S.

Name(s): Phu Dang

Website Link: <https://pndang.com/illuminating-cognizance/>

Code

```
In [ ]: # Importing packages and libraries

import pandas as pd
import numpy as np
import os
from scipy.stats import ks_2samp

import plotly.express as px
import plotly.graph_objects as go
pd.options.plotting.backend = 'plotly'

In [ ]: # Adjust dataframe display options to view full output
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)
# pd.set_option('display.max_rows', None)

# Reset to default, comment-out 2 (or 3) lines above and
# uncomment 2 lines below, run all before pushing to GitHub
# Reset to default helps limit the display scope of output dataframes,
# easier to navigate notebook
# pd.set_option('display.max_columns', 20)
# pd.set_option('display.width', 80)
```

Introduction

Analysis questions of interest:

1. Are there any possible connections between outage duration and regional consumption information?
 - Look into different types of information: price vs. consumption vs. customers served
 - **(Personal favorite)** I plan to investigate this thoroughly
2. Is there a relation between anomaly level and cause category?

- If there are trends, are they statistically significant?
- Do extreme anomaly levels appear to be related to different cause categories than regular levels?
- **(New Personal favorite)** I plan to investigate this thoroughly

3. Do the start times of outages seem to affect outage duration?

- Rationale: A sudden outage at night can give workers more time and space to fix overnight

4. What are some common characteristics of longer outages?

Cleaning and EDA

In []: *# Importing dataset*

```
path = os.path.join('data', 'outage.csv').replace('\\', '/')

df = pd.read_csv(path, header=5, skiprows=[6])
df = df.loc[:, ~df.columns.isin(['OBS', 'variables'])]
df.head()
```

Out[]:

	YEAR	MONTH	U.S.STATE	POSTAL.CODE	NERC.REGION	CLIMATE.REGION	ANOMALY.
0	2011	7.0	Minnesota	MN	MRO	East North Central	
1	2014	5.0	Minnesota	MN	MRO	East North Central	
2	2010	10.0	Minnesota	MN	MRO	East North Central	
3	2012	6.0	Minnesota	MN	MRO	East North Central	
4	2015	7.0	Minnesota	MN	MRO	East North Central	

In []: `df.shape`

Out[]: (1534, 55)

In []: `df.columns`

```
Out[ ]: Index(['YEAR', 'MONTH', 'U.S._STATE', 'POSTAL.CODE', 'NERC.REGION',
             'CLIMATE.REGION', 'ANOMALY.LEVEL', 'CLIMATE.CATEGORY',
             'OUTAGE.START.DATE', 'OUTAGE.START.TIME', 'OUTAGE.RESTORATION.DATE',
             'OUTAGE.RESTORATION.TIME', 'CAUSE.CATEGORY', 'CAUSE.CATEGORY.DETAIL',
             'HURRICANE.NAMES', 'OUTAGE.DURATION', 'DEMAND.LOSS.MW',
             'CUSTOMERS.AFFECTED', 'RES.PRICE', 'COM.PRICE', 'IND.PRICE',
             'TOTAL.PRICE', 'RES.SALES', 'COM.SALES', 'IND.SALES', 'TOTAL.SALES',
             'RES.PERCEN', 'COM.PERCEN', 'IND.PERCEN', 'RES.CUSTOMERS',
             'COM.CUSTOMERS', 'IND.CUSTOMERS', 'TOTAL.CUSTOMERS', 'RES.CUST.PCT',
             'COM.CUST.PCT', 'IND.CUST.PCT', 'PC.REALGSP.STATE', 'PC.REALGSP.USA',
             'PC.REALGSP.REL', 'PC.REALGSP.CHANGE', 'UTIL.REALGSP', 'TOTAL.REALGSP',
             'UTIL.CONTRI', 'PI.UTIL.OFUSA', 'POPULATION', 'POPPCT_URBAN',
             'POPPCT_UC', 'POPDEN_URBAN', 'POPDEN_UC', 'POPDEN_RURAL',
             'AREAPCT_URBAN', 'AREAPCT_UC', 'PCT_LAND', 'PCT_WATER_TOT',
             'PCT_WATER_INLAND'],
            dtype='object')
```

```
In [ ]: df.head()
```

```
Out[ ]:   YEAR  MONTH  U.S._STATE  POSTAL.CODE  NERC.REGION  CLIMATE.REGION  ANOMALY.
```

	YEAR	MONTH	U.S._STATE	POSTAL.CODE	NERC.REGION	CLIMATE.REGION	ANOMALY.
0	2011	7.0	Minnesota	MN	MRO	East North Central	
1	2014	5.0	Minnesota	MN	MRO	East North Central	
2	2010	10.0	Minnesota	MN	MRO	East North Central	
3	2012	6.0	Minnesota	MN	MRO	East North Central	
4	2015	7.0	Minnesota	MN	MRO	East North Central	

```
In [ ]: df.dtypes
```

```
Out[ ]: YEAR                int64
MONTH                float64
U.S._STATE            object
POSTAL.CODE           object
NERC.REGION           object
...
AREAPCT_URBAN         float64
AREAPCT_UC            float64
PCT_LAND              float64
PCT_WATER_TOT         float64
PCT_WATER_INLAND     float64
Length: 55, dtype: object
```

```
In [ ]: # Combining OUTAGE.START.DATE and OUTAGE.START.TIME

dates_start = df['OUTAGE.START.DATE'] + ' ' + df['OUTAGE.START.TIME']
df['OUTAGE.START'] = pd.to_datetime(dates_start)
```

```

dates_end = df['OUTAGE.RESTORATION.DATE'] + ' ' + df['OUTAGE.RESTORATION.TIME']
df['OUTAGE.END'] = pd.to_datetime(dates_end)

df.drop(columns=['OUTAGE.START.DATE', \
                'OUTAGE.RESTORATION.DATE', 'OUTAGE.RESTORATION.TIME'], inplace=True)

```

```
In [ ]: # Get counts of missing values in each column
```

```
df.isna().sum()
```

```

Out[ ]: YEAR                0
        MONTH              9
        U.S._STATE         0
        POSTAL.CODE        0
        NERC.REGION        0
        ..
        PCT_LAND           0
        PCT_WATER_TOT      0
        PCT_WATER_INLAND   0
        OUTAGE.START       9
        OUTAGE.END        58
        Length: 54, dtype: int64

```

```
In [ ]: df.dtypes
```

```

Out[ ]: YEAR                int64
        MONTH              float64
        U.S._STATE         object
        POSTAL.CODE        object
        NERC.REGION        object
        ...
        PCT_LAND           float64
        PCT_WATER_TOT      float64
        PCT_WATER_INLAND   float64
        OUTAGE.START       datetime64[ns]
        OUTAGE.END         datetime64[ns]
        Length: 54, dtype: object

```

```
In [ ]: # Adding a duration column in hours
```

```
df['DURATION.HR'] = df['OUTAGE.DURATION'] / 60
```

```
In [ ]: # Check min and max duration
```

```

print(df['DURATION.HR'].min())
print(df['DURATION.HR'].max())

```

```

0.0
1810.8833333333334

```

```
In [ ]: df.head()
```

Out []:

	YEAR	MONTH	U.S._STATE	POSTAL.CODE	NERC.REGION	CLIMATE.REGION	ANOMALY.
0	2011	7.0	Minnesota	MN	MRO	East North Central	
1	2014	5.0	Minnesota	MN	MRO	East North Central	
2	2010	10.0	Minnesota	MN	MRO	East North Central	
3	2012	6.0	Minnesota	MN	MRO	East North Central	
4	2015	7.0	Minnesota	MN	MRO	East North Central	

In []:

```
# Convert dataframe to markdown

print(df[['U.S._STATE', 'POSTAL.CODE', 'ANOMALY.LEVEL', 'OUTAGE.START.TIME',
          'CAUSE.CATEGORY', 'CAUSE.CATEGORY.DETAIL', 'DEMAND.LOSS.MW', 'RES.PRICE',
          'PC.REALGSP.STATE', 'PCT_WATER_INLAND', 'DURATION.HR']].head().to_markdown(index=False, na_rep='nan'))
```

U.S._STATE	POSTAL.CODE	ANOMALY.LEVEL	OUTAGE.START.TIME	CAUSE.CATEGORY	CAUSE.CATEGORY.DETAIL	DEMAND.LOSS.MW	RES.PRICE	PC.REALGSP.STATE
Minnesota	MN	-0.3	5:00:00 PM	severe weather	nan	11.6	51	268
Minnesota	MN	-0.1	6:38:00 PM	intentional attack	vandalism	nan	12.12	534
Minnesota	MN	-1.5	8:00:00 PM	severe weather	heavy wind	nan	10.87	50
Minnesota	MN	-0.1	4:30:00 AM	severe weather	thunderstorm	nan	11.79	51
Minnesota	MN	1.2	2:00:00 AM	severe weather	nan	250	13.07	54

In []:

```
# Change postal code column name to STATE.ABBR (state abbreviation)

df.rename(columns={'POSTAL.CODE': 'STATE.ABBR'}, inplace=True)
```

Univariate Analysis

In []:

```
# Plotting the distribution of outage durations
```

```
fig1 = px.histogram(df, x='DURATION.HR', nbins=200, histnorm='probability', \
                    title='Distribution of outage duration, in hours')
fig1.update_layout(xaxis_range=[0, 360])
fig1.update_xaxes(title_text='duration in hours')
fig1
```

In []: *# Write to html file*

```
path = os.path.join('assets', 'uni_1.html')
fig1.write_html(path, include_plotlyjs='cdn')
```

In []: *# Sanity check*

```
573 / len(df['DURATION.HR'])
```

Out[]: 0.37353324641460234

In []: *# Plotting the distribution of cause category*

```
fig2 = px.histogram(df, x='CAUSE.CATEGORY', histnorm='probability', \
                    title='Distribution of outage cause category')
fig2.update_xaxes(title_text='cause category')
fig2
```

In []: *# Write to html file*

```
path = os.path.join('assets', 'uni_2.html')
fig2.write_html(path, include_plotlyjs='cdn')
```

Bivariate Analysis

In []: *# Plotting a scatterplot between outage duration and residential electricity price*

```
fig3 = px.scatter(df, x='RES.PRICE', y='DURATION.HR', title='Outage duration and re
fig3.update_xaxes(title_text='cents/kilowatt-hour')
fig3.update_yaxes(title_text='hours')
fig3
```

In []: *# Write to html file*

```
path = os.path.join('assets', 'bi_1.html')
fig3.write_html(path, include_plotlyjs='cdn')
```

In []: *# Plotting a scatterplot between outage duration and customers served*

```
fig4 = px.scatter(df, x='TOTAL.CUSTOMERS', y='DURATION.HR' \
                  , title='Outage duration and total number of customers served, annually')
fig4.update_xaxes(title_text='number of customers served')
fig4.update_yaxes(title_text='hours')
fig4
```

```
In [ ]: # Plotting a barplot between outage duration and regional economic output

fig5 = px.scatter(df, x='PC.REALGSP.STATE', y='DURATION.HR'\
, title='Outage duration and per capita GSP')
fig5.update_xaxes(title_text='per capita real gross state product (measured in 2009
fig5.update_yaxes(title_text='hours')
fig5
```

```
In [ ]: # Write to html file

path = os.path.join('assets', 'bi_2.html')
fig5.write_html(path, include_plotlyjs='cdn')
```

Interesting Aggregates

```
In [ ]: # Plot average outage duration for every combination of state and cause category

group1 = pd.pivot_table(df, index=['U.S._STATE'], columns=['CAUSE.CATEGORY'], \
values='DURATION.HR', aggfunc=np.mean)
group1['dummy'] = group1.sum(axis=1)
group1 = group1.sort_values(by='dummy', ascending=True).drop(columns=['dummy'])
fig6 = group1.plot(kind='barh', \
title='Average outage duration by state and cause category',
labels={'U.S._STATE': 'state ', 'value': 'average outage duration in hours ', \
'CAUSE.CATEGORY': 'cause category '})
fig6.update_layout(height=1000, width=1500, legend=dict(title='cause category'))
fig6
```

```
In [ ]: # Write to html file

path = os.path.join('assets', 'multi_1.html')
fig6.write_html(path, include_plotlyjs='cdn')
```

```
In [ ]: # Sanity check

group1.sum(axis=1).sort_values(ascending=False).index
```

```
Out[ ]: Index(['Michigan', 'Louisiana', 'Wisconsin', 'New York', 'Arizona', 'Indiana',
'Texas', 'Kentucky', 'California', 'Florida', 'Kansas', 'West Virginia',
'Iowa', 'Washington', 'New Jersey', 'Pennsylvania', 'Ohio', 'Illinois',
'Oklahoma', 'Missouri', 'District of Columbia', 'Massachusetts',
'Tennessee', 'Maryland', 'Arkansas', 'Maine', 'Minnesota', 'Utah',
'Nebraska', 'South Carolina', 'Colorado', 'Oregon', 'North Carolina',
'Connecticut', 'Delaware', 'Virginia', 'Idaho', 'New Hampshire',
'Georgia', 'Alabama', 'Hawaii', 'North Dakota', 'Nevada', 'Mississippi',
'New Mexico', 'Wyoming', 'Montana', 'South Dakota', 'Vermont'],
dtype='object', name='U.S._STATE')
```

```
In [ ]: group1
```

Out[]:

CAUSE.CATEGORY	equipment failure	fuel supply emergency	intentional attack	islanding	public appeal	severe weather	order
U.S._STATE							
Vermont	NaN	NaN	0.590741	NaN	NaN	NaN	
South Dakota	NaN	NaN	NaN	2.000000	NaN	NaN	
Montana	NaN	NaN	1.550000	0.575000	NaN	NaN	
Wyoming	1.016667	NaN	0.005556	0.533333	NaN	1.766667	
New Mexico	NaN	1.266667	2.908333	NaN	NaN	NaN	
...	
Arizona	2.308333	NaN	10.660000	NaN	NaN	428.775000	
New York	4.116667	278.120833	5.151389	NaN	44.250000	100.576263	
Wisconsin	NaN	566.187500	7.650000	NaN	6.466667	25.457143	
Louisiana	2.938889	469.500000	NaN	NaN	22.653571	119.782143	
Michigan	440.588889	NaN	60.587500	0.016667	17.966667	80.527510	

49 rows × 7 columns

In []:

```
# Get aggregated mean outage duration in minutes, hours, and customers affected by
# pd.set_option('display.max_rows', None)
group2 = df.groupby(by=['STATE.ABBR', 'CAUSE.CATEGORY']).mean()[['OUTAGE.DURATION',
group2.rename(columns={'OUTAGE.DURATION': 'Avg duration (mins)', \
'DURATION.HR': 'Avg duration (hrs)', \
'CUSTOMERS.AFFECTED': 'Avg # of customers affected'})
pd.set_option('display.max_rows', 20)
group2.head()
```

Out[]:

		OUTAGE.DURATION	DURATION.HR	CUSTOMERS.AFFECTED
STATE.ABBR	CAUSE.CATEGORY			
AK	equipment failure	NaN	NaN	14273.0
AL	intentional attack	77.000000	1.283333	NaN
	severe weather	1421.750000	23.695833	94328.8
AR	equipment failure	105.000000	1.750000	NaN
	intentional attack	547.833333	9.130556	9200.0


```
In [ ]: # sanity check
```

```
df[df['STATE.ABBR'] == 'AK']
```

```
Out[ ]:
```

	YEAR	MONTH	U.S._STATE	STATE.ABBR	NERC.REGION	CLIMATE.REGION	ANOMALY
1533	2000	NaN	Alaska	AK	ASCC	NaN	

```
In [ ]: print(group1[-10:-5].to_markdown(index=True))
```

U.S._STATE	equipment failure	fuel supply emergency	intentional attack
islanding	public appeal	severe weather	system operability disruption
Florida	9.24167	nan	0.83333
nan	72	107.003	3.428
California	8.74683	102.577	15.7743
3.58095	33.8019	48.8062	6.06111
Kentucky	10.8667	209.5	1.8
nan	nan	74.6685	nan
Texas	6.76	232	4.97949
nan	19.0069	64.2482	13.5133
Indiana	0.0166667	204	7.03125
2.08889	nan	75.3882	77.86

```
In [ ]: # Get pivot table of average anomaly level by cause category and category detail
```

```
group2 = pd.pivot_table(df, index=['CAUSE.CATEGORY'], \
    columns=['CAUSE.CATEGORY.DETAIL'], values='ANOMALY.LEVEL', aggfunc=np.mean)
fig7 = group2.plot(kind='barh', \
    title='Average anomaly level by cause category, subset by cause '+
    'detail<br><sup>Anomaly level represents the oceanic El Niño/La Niña (ONI)<sup>'+
    ' index, estimated as a 3-month running mean of ERSST.v4 SST anomalies in'+
    ' the Niño 3.4 region</sup>',
    labels={'CAUSE.CATEGORY': 'cause category', 'value': 'average anomaly level (ON
    'CAUSE.CATEGORY.DETAIL': 'cause detail'})
fig7.update_layout(legend=dict(title='cause detail'))
fig7
```

```
In [ ]: # Write to html file
```

```
path = os.path.join('assets', 'oni_by_cause.html')
fig7.write_html(path, include_plotlyjs='cdn')
```

```
In [ ]: group2
```

Out []:

CAUSE.CATEGORY.DETAIL	Coal	Hydro	Natural Gas	100 MW loadshed	Coal	HVSubstation interruption	Hyd
CAUSE.CATEGORY							
equipment failure	NaN	NaN	NaN	NaN	NaN	NaN	NaN
fuel supply emergency	-0.29	-0.4	-0.514286	NaN	-0.471429	NaN	-0.4
intentional attack	NaN	NaN	NaN	NaN	NaN	NaN	NaN
severe weather	NaN	NaN	NaN	NaN	NaN	NaN	NaN
system operability disruption	NaN	NaN	NaN	-0.2	NaN	-0.1	NaN

In []:

```
# Write to html file

path = os.path.join('assets', 'multi_2.html')
fig7.write_html(path, include_plotlyjs='cdn')
```

In []:

```
df.head()
```

Out []:

	YEAR	MONTH	U.S.STATE	STATE.ABBR	NERC.REGION	CLIMATE.REGION	ANOMALY.LE
0	2011	7.0	Minnesota	MN	MRO	East North Central	
1	2014	5.0	Minnesota	MN	MRO	East North Central	
2	2010	10.0	Minnesota	MN	MRO	East North Central	
3	2012	6.0	Minnesota	MN	MRO	East North Central	
4	2015	7.0	Minnesota	MN	MRO	East North Central	

Assessment of Missingness

In []:

```
# Inspect the outages where month is missing (likely MAR, explained by year)

df[df['MONTH'].isna()].head(11)
```

Out[]:

	YEAR	MONTH	U.S._STATE	STATE.ABBR	NERC.REGION	CLIMATE.REGION	ANOMAL
239	2000	NaN	Texas	TX	FRCC	South	
339	2000	NaN	Alabama	AL	SERC	Southeast	
365	2000	NaN	Illinois	IL	SERC	Central	
766	2000	NaN	North Carolina	NC	SERC	Southeast	
887	2000	NaN	Delaware	DE	RFC	Northeast	
1318	2000	NaN	Virginia	VA	SERC	Southeast	
1506	2002	NaN	Kansas	KS	SPP	South	
1530	2006	NaN	North Dakota	ND	MRO	West North Central	
1533	2000	NaN	Alaska	AK	ASCC	NaN	

In []:

```
# Inspect the outages where peak-hours demand loss is missing (demand loss is likely  
df[df['DEMAND_LOSS_MW'].isna()].head(7)
```

Out[]:

	YEAR	MONTH	U.S._STATE	STATE.ABBR	NERC.REGION	CLIMATE.REGION	ANOMALY.LE
0	2011	7.0	Minnesota	MN	MRO	East North Central	
1	2014	5.0	Minnesota	MN	MRO	East North Central	
2	2010	10.0	Minnesota	MN	MRO	East North Central	
3	2012	6.0	Minnesota	MN	MRO	East North Central	
5	2010	11.0	Minnesota	MN	MRO	East North Central	
6	2010	7.0	Minnesota	MN	MRO	East North Central	
9	2013	6.0	Minnesota	MN	MRO	East North Central	

Missingness Assessment Analysis notes:

DEMAND.LOSS.MW could possibly **depend** on outage start time and/or outage duration

- Rationale: A short outage outside of high-demand times (4 PM - 9 PM) will not have data for DEMAND.LOSS.MW

DEMAND.LOSS.MW could also **depend** on the number of customers affected (anticipating a positive correlation)

DEMAND.LOSS.MW is likely to **not depend** on PCT_WATER_INLAND

- PCT_WATER_INLAND ~ percentage of inland water area in the U.S. state as compared to the overall inland water area in the continental U.S. (in %)

```
In [ ]: # Query out relevant columns for assessing missingness to keep dataframe simple

aom = df[['DEMAND.LOSS.MW', 'OUTAGE.START.TIME', 'DURATION.HR', \
          'CUSTOMERS.AFFECTED', 'PCT_WATER_INLAND']]
print(aom.shape[0])
aom.head()
```

1534

```
Out [ ]: DEMAND.LOSS.MW  OUTAGE.START.TIME  DURATION.HR  CUSTOMERS.AFFECTED  PCT_W
```

	DEMAND.LOSS.MW	OUTAGE.START.TIME	DURATION.HR	CUSTOMERS.AFFECTED	PCT_W
0	NaN	5:00:00 PM	51.000000	70000.0	
1	NaN	6:38:00 PM	0.016667	NaN	
2	NaN	8:00:00 PM	50.000000	70000.0	
3	NaN	4:30:00 AM	42.500000	68200.0	
4	250.0	2:00:00 AM	29.000000	250000.0	

```
In [ ]: # Inspect the missingness of columns in aom

aom.isna().sum()
```

```
Out [ ]: DEMAND.LOSS.MW      705
OUTAGE.START.TIME         9
DURATION.HR              58
CUSTOMERS.AFFECTED      443
PCT_WATER_INLAND         0
dtype: int64
```

Permutation test wrt outage duration

- Assess missingness of DEMAND.LOSS.MW

```
In [ ]: # Assign a boolean column DEMAND.MISSING that indicates whether, or not, peak
# demand lost is missing
```

```
aom['DEMAND.MISSING'] = aom['DEMAND.LOSS.MW'].isna()
aom.head(3)
```

C:\Users\phuro\AppData\Local\Temp\ipykernel_3788\1607562000.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
Out[ ]: DEMAND.LOSS.MW  OUTAGE.START.TIME  DURATION.HR  CUSTOMERS.AFFECTED  PCT_W
```

0	NaN	5:00:00 PM	51.000000	70000.0
1	NaN	6:38:00 PM	0.016667	NaN
2	NaN	8:00:00 PM	50.000000	70000.0

```
In [ ]: # Box plot of duration hours by missingness of peak-hours demand Loss
```

```
aom_fig1 = px.box(aom, x='DURATION.HR', color='DEMAND.MISSING')
aom_fig1.update_layout(xaxis_range=[0, 500])
aom_fig1
```

```
In [ ]: # Histogram of duration hours by missingness of peak-hours demand Loss
```

```
aom_fig2 = px.histogram(aom, x='DURATION.HR', color='DEMAND.MISSING', \
    histnorm='probability', nbins=3000, marginal='rug')
aom_fig2.update_layout(xaxis_range=[0, 24])
aom_fig2
```

```
In [ ]: aom.dtypes
```

```
Out[ ]: DEMAND.LOSS.MW      float64
OUTAGE.START.TIME        object
DURATION.HR              float64
CUSTOMERS.AFFECTED      float64
PCT_WATER_INLAND        float64
DEMAND.MISSING           bool
dtype: object
```

```
In [ ]: # Convert outage start times to datetime data type
```

```
times = pd.to_datetime(aom['OUTAGE.START.TIME'])
aom['START.TIME'] = times
aom = aom.sort_values(by='START.TIME', ascending=True)
aom.head(5)
```

C:\Users\phuro\AppData\Local\Temp\ipykernel_3788\3097801354.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
Out[ ]:
```

	DEMAND.LOSS.MW	OUTAGE.START.TIME	DURATION.HR	CUSTOMERS.AFFECTED	PC
557	0.0	12:00:00 AM	0.016667	0.0	
545	2000.0	12:00:00 AM	214.833333	650000.0	
544	0.0	12:00:00 AM	0.033333	0.0	
1374	10.0	12:00:00 AM	13.500000	47165.0	
542	0.0	12:00:00 AM	24.000000	0.0	

```
In [ ]: # Histogram of outage start time by missingness of peak-hours demand lost

mean = aom['START.TIME'].mean()
median = aom['START.TIME'].median()

aom_fig3 = px.histogram(aom, x='START.TIME', color='DEMAND.MISSING', \
    histnorm='probability',
    title='Outage start time by missingness of peak demand lost',
    labels={'START.TIME': 'outage start time', 'DEMAND.MISSING': 'demand lost missi
    barmode='overlay', marginal='box', opacity=0.7)

aom_fig3.add_vline(x=median, line_color='#00CC96')
aom_fig3.add_vline(x=mean, line_color='#AB63FA')

aom_fig3.update_xaxes(
    ticktext=["00:00", "03:00", "06:00", "09:00", "12:00", "14:30", "18:00", \
        "21:00<br><b>|<b><br>", "00:00", "16:00<br><b>|<b><br>",
        "<br>peak demand window<br>", "|<br>mean<br>",
        "|<br>|<br>median<br><br>"],
    tickvals=[pd.to_datetime('00:00:00'), pd.to_datetime('03:00:00'), \
        pd.to_datetime('06:00:00'), pd.to_datetime('09:00:00'), \
        pd.to_datetime('12:00:00'), pd.to_datetime('14:30:00'),
        pd.to_datetime('18:00:00'), pd.to_datetime('21:00:00'),
        pd.to_datetime('23:59:00'), pd.to_datetime('16:00:00'),
        pd.to_datetime('18:30:00'), mean, median],
    tickangle=0
)
aom_fig3
```

```
In [ ]: # Write to html file
```

```
path = os.path.join('assets', 'aom_perm1_observed_dist.html')
aom_fig3.write_html(path, include_plotlyjs='cdn')
```

```
In [ ]: aom.head(3)
```

```
Out[ ]:
```

	DEMAND.LOSS.MW	OUTAGE.START.TIME	DURATION.HR	CUSTOMERS.AFFECTED	PCT
557	0.0	12:00:00 AM	0.016667	0.0	
545	2000.0	12:00:00 AM	214.833333	650000.0	
544	0.0	12:00:00 AM	0.033333	0.0	

Null Hypothesis: The start times between outages where the amount of peak demand lost is missing, and outages where the amount of peak demand lost is **not** missing, have **the same** distribution. Any observed difference is due to chance alone.

Alternative Hypothesis: The outage start times by missingness of peak demand lost have **different** distributions. The observed difference is **unlikely** due to chance alone.

Test statistic: difference in group median start time

Significance level: 5%

Method: shuffle DEMAND.MISSING (status of missing) column to simulate under null hypothesis

```
In [ ]: # Run simulation 5000 times
```

```
N = 5000
differences = []

for _ in range(N):
    aom['shuffled'] = np.random.permutation(aom['DEMAND.MISSING'])
    w_missing_median = aom[aom['shuffled']]['START.TIME'].median()
    wo_missing_median = aom[~aom['shuffled']]['START.TIME'].median()
    test_stat = abs(w_missing_median - wo_missing_median)
    differences.append(test_stat)
```

```
In [ ]: # Get observed test statistic
```

```
w_missing_median_obs = aom[aom['DEMAND.MISSING']]['START.TIME'].median()
wo_missing_median_obs = aom[~aom['DEMAND.MISSING']]['START.TIME'].median()
observed = abs(w_missing_median_obs - wo_missing_median_obs)
observed = observed.total_seconds()
```

```
In [ ]: # Get p-value
```

```
results = []
differences_secs = []
```

```
for val in differences:
    val = val.total_seconds()
    differences_secs.append(val)
    if val >= observed:
        results.append(True)
        continue
    results.append(False)

pval = np.mean(results)
pval
```

Out[]: 0.1796

```
In [ ]: # Plot distribution of results

perm1 = px.histogram(pd.DataFrame({'simulated differences': differences_secs}),\
    x='simulated differences', histnorm='probability',
    title='Simulated differences of median outage start times',
    # <br><sup>Calculation: absolute difference between median outage start tim
    labels={'simulated differences': 'simulated differences (seconds)'})
perm1.add_vline(x=observed, line_color='rgb(0,100,80)',
    annotation_text='observed', annotation_position='top left')
p_95 = np.percentile(differences_secs, 95)
perm1.add_vline(x=p_95, line_color='rgb(0,176,246)',
    annotation_text='significance level (5%)', annotation_position='top rig
perm1
```

```
In [ ]: # Write to html file

path = os.path.join('assets', 'aom_perm1_results.html')
perm1.write_html(path, include_plotlyjs='cdn')
```

Result: **Fail to reject** the null hypothesis at a 5% significance level

- Missingness of peak demand lost is likely to not depend on outage start times

```
In [ ]: aom.head(3)
```

Out[]:

	DEMAND.LOSS.MW	OUTAGE.START.TIME	DURATION.HR	CUSTOMERS.AFFECTED	PCT.
557	0.0	12:00:00 AM	0.016667		0.0
545	2000.0	12:00:00 AM	214.833333	650000.0	
544	0.0	12:00:00 AM	0.033333		0.0

PCT_WATER_INLAND

- Percentage of inland water area in the U.S. state as compared to the overall inland water area in the continental U.S. (in %)


```

In [ ]: # Histogram of state water percentage by missingness of peak-hours demand loss

aom_fig4 = px.histogram(aom, x='PCT_WATER_INLAND', color='DEMAND.MISSING', \
    histnorm='probability', barmode='overlay', marginal='box', opacity=0.7,
    nbins=20, title="State water percentage by missingness of peak demand lost",
    labels={'PCT_WATER_INLAND': 'state water percentage',
    'DEMAND.MISSING': 'demand lost missing'})

aom_fig4.add_vline(x=aom['PCT_WATER_INLAND'].mean(), annotation_text='mean', \
    line_color='#AB63FA')
aom_fig4.add_vline(x=aom['PCT_WATER_INLAND'].median(), annotation_text='median', \
    line_color='#00CC96')
aom_fig4.update_layout(yaxis_range=[0, 0.37])
aom_fig4

In [ ]: # Write to html file

path = os.path.join('assets', 'aom_perm2_observed_dist.html')
aom_fig4.write_html(path, include_plotlyjs='cdn')

In [ ]: # Plot the CDFs of state water percentage by missingness of peak demand lost

def create_cdf(df, group_col, group, stat_col):
    return df.loc[df[group_col] == group, stat_col].value_counts(normalize=True).sort_index()

cdf_fig = go.Figure()

cdf_fig.add_trace(
    go.Scatter(x=create_cdf(aom, 'DEMAND.MISSING', True, 'PCT_WATER_INLAND').index,
    y=create_cdf(aom, 'DEMAND.MISSING', True, 'PCT_WATER_INLAND'),
    name='CDF with missing demand')
)

cdf_fig.add_trace(
    go.Scatter(x=create_cdf(aom, 'DEMAND.MISSING', False, 'PCT_WATER_INLAND').index,
    y=create_cdf(aom, 'DEMAND.MISSING', False, 'PCT_WATER_INLAND'),
    name='CDF without missing demand')
)

cdf_fig.update_layout(title='CDFs of state water percentage by missingness of peak
cdf_fig.update_xaxes(title='state water percentage')
cdf_fig.update_yaxes(title='cumulative probability')

In [ ]: # Write to html file

path = os.path.join('assets', 'aom_perm2_cdf.html')
cdf_fig.write_html(path, include_plotlyjs='cdn')

```

Null Hypothesis: The distribution of state inland water percentage in outages where peak demand lost **is** missing is **the same** as in outages where peak demand lost **is not** missing. Any observed difference is due to chance alone.

Alternative Hypothesis: The distributions of state inland water percentage between the two groups are different. The observed difference is **unlikely** due to chance alone.

Test statistic: K-S statistic

Significance level: 5%

Method: shuffle DEMAND.MISSING (status of missing) column to simulate under null hypothesis

```
In [ ]: # Run simulation 5000 times

N = 10000
results = []

for _ in range(N):
    aom['shuffled'] = np.random.permutation(aom['DEMAND.MISSING'])
    groups = aom.groupby('shuffled')['PCT_WATER_INLAND']
    ks_stat = ks_2samp(groups.get_group(True), groups.get_group(False)).statistic
    results.append(ks_stat)

results[:5]
```

```
Out[ ]: [0.02988305144196631,
         0.038723917562816006,
         0.046476571790330996,
         0.050458126940943974,
         0.03833038181522641]
```

```
In [ ]: # Get observed K-S statistic

observed_ks = ks_2samp(aom.loc[aom['DEMAND.MISSING'], 'PCT_WATER_INLAND'], \
    aom.loc[~aom['DEMAND.MISSING'], 'PCT_WATER_INLAND']).statistic
observed_ks
```

```
Out[ ]: 0.08352539588840695
```

```
In [ ]: # Get p-value

pval = (np.array(results) >= observed_ks).mean()
pval
```

```
Out[ ]: 0.0036
```

```
In [ ]: # Plot simulated K-S stats

perm2 = px.histogram(pd.DataFrame(data={'simulated K-S statistics': results}), \
    x='simulated K-S statistics', histnorm='probability',
    title='Empirical distribution of the K-S Statistic, state water percentage by m
perm2.add_vline(x=observed_ks, line_color='rgb(0,100,80)', \
    annotation_text='observed K-S')
p_95 = np.percentile(results, 95)
perm2.add_vline(x=p_95, line_color='rgb(0,176,246)', \
```

```
annotation_text='significance level (5%)', annotation_position='top left')
perm2
```

```
In [ ]: # Write to html file

path = os.path.join('assets', 'aom_perm2_results.html')
perm2.write_html(path, include_plotlyjs='cdn')
```

Result: **Reject** the null hypothesis at a 5% significance level

- Missingness of peak demand lost **could possibly** depend on the state proportion of inland water relative to continental U.S. (MAR dependent)

```
In [ ]: # sanity check

aom['PCT_WATER_INLAND'].min()
```

Out[]: 0.240151328

Hypothesis Testing

Test the observed distribution of anomaly level by cause category

```
In [ ]: df['ANOMALY.LEVEL'].plot()
```

```
In [ ]: fig7
```

```
In [ ]: group2
```

Out[]:

CAUSE.CATEGORY.DETAIL	Coal	Hydro	Natural Gas	100 MW loadshed	Coal	HVSubstation interruption	Hyd
CAUSE.CATEGORY							
equipment failure	NaN	NaN	NaN	NaN	NaN	NaN	NaN
fuel supply emergency	-0.29	-0.4	-0.514286	NaN	-0.471429	NaN	-0.
intentional attack	NaN	NaN	NaN	NaN	NaN	NaN	NaN
severe weather	NaN	NaN	NaN	NaN	NaN	NaN	NaN
system operability disruption	NaN	NaN	NaN	-0.2	NaN	-0.1	NaN

```
In [ ]: group2.abs().sum(axis=1).plot(kind='bar')
```

```
In [ ]: # Inspect the distribution of anomaly levels (observed)
```

```

oni_fig = df[['ANOMALY.LEVEL', 'CAUSE.CATEGORY']].plot(kind='hist', \
x='ANOMALY.LEVEL', histnorm='probability',
labels={'ANOMALY.LEVEL': 'anomaly level (ONI Index)', 'CAUSE.CATEGORY':
'cause category'},
title='Distribution of anomaly level (ONI Index), subset by cause category',
color='CAUSE.CATEGORY')
oni_fig.add_vline(x=df['ANOMALY.LEVEL'].mean(), annotation_text='mean')
oni_fig.add_vline(x=df['ANOMALY.LEVEL'].median(), annotation_text='median', \
annotation_position='top left')
oni_fig.update_layout(yaxis_range=[0, 0.9])

```

In []: *# Write to html file*

```

path = os.path.join('assets', 'hyp1_oni_dist.html')
oni_fig.write_html(path, include_plotlyjs='cdn')

```

In []: *# Inspect the distribution of anomaly levels (observed) separately by cause categor*

```

for cat in group2.index:
    data = df[df['CAUSE.CATEGORY'] == cat]
    plot = data['ANOMALY.LEVEL'].plot(kind='hist', histnorm='probability', labels={
'value': 'anomaly level (ONI Index)'},
title=f'Distribution of anomaly levels for {cat}')
    plot.add_vline(x=data['ANOMALY.LEVEL'].mean(), annotation_text='mean')
    plot.add_vline(x=data['ANOMALY.LEVEL'].median(), annotation_text='median', \
annotation_position='top left')
    plot.show()

# Write to html file
path = os.path.join('assets', f'hyp1_oni_obs_dist_{cat}.html')
plot.write_html(path, include_plotlyjs='cdn')

```

Conclusion from above: A non-parametric test (permutation) seems appropriate because the observed distribution of anomaly level is non-normal

Test begin

Null Hypothesis: The distribution of outage cause categories with **extreme** anomaly levels is **the same** as in outage cause categories with **regular** anomaly levels. Any observed difference is due to chance alone.

Alternative Hypothesis: The distributions of outage cause categories between the two groups are different. The observed difference is **unlikely** due to chance alone.

Test statistic: TVD

Significance level: 5%

Method: shuffle the binary column that indicates whether the associated anomaly level is extreme to simulate under null hypothesis

- Anomaly levels above 0.5 or below -0.5 are considered extreme and are indicative of El Nino and La Nina, respectively

```
In [ ]: def get_tvd(df, groups, cats, show_steps=False):
    """
        Function to calculate the Total Variation Distance between groups

        groups: the binary column
        cats: the categorical column
    """

    # Get count of each pair of category and group
    counts = df.pivot_table(index=cats, columns=groups, aggfunc='size')

    # Normalize each column
    distr = counts / counts.sum()

    if show_steps:
        print("\nSTEPS:")
        print(counts)
        print()
        print(distr)
        print()
        print(distr.diff(axis=1))
        print()
        print(distr.diff(axis=1).iloc[:, -1])
        print()
        print(distr.diff(axis=1).iloc[:, -1].abs())
        print()
        print(distr.diff(axis=1).iloc[:, -1].abs().sum() / 2)
        print("END\n")

    # Compute and return the total variation distance
    return distr.diff(axis=1).iloc[:, -1].abs().sum() / 2
```

```
In [ ]: # Assign an "extreme" column that indicates whether, or not, the associated
# anomaly level is extreme

oni = df[['ANOMALY.LEVEL', 'CAUSE.CATEGORY']]
oni['extreme'] = (oni['ANOMALY.LEVEL'] < -0.5) | (oni['ANOMALY.LEVEL'] > 0.5)
oni.head(10)
```

C:\Users\phuro\AppData\Local\Temp\ipykernel_3788\45707634.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[]:

	ANOMALY.LEVEL	CAUSE.CATEGORY	extreme
0	-0.3	severe weather	False
1	-0.1	intentional attack	False
2	-1.5	severe weather	True
3	-0.1	severe weather	False
4	1.2	severe weather	True
5	-1.4	severe weather	True
6	-0.9	severe weather	True
7	0.2	severe weather	False
8	0.6	intentional attack	True
9	-0.2	severe weather	False

```
In [ ]: # Get observed test statistic

observed_tvd = get_tvd(oni, groups='extreme', cats='CAUSE.CATEGORY')
observed_tvd
```

Out[]: 0.09399855386840203

```
In [ ]: # Run simulation 5000

N = 5000
results = []

for _ in range(N):
    oni['shuffled'] = np.random.permutation(oni['extreme'])
    tvd = get_tvd(oni, groups='shuffled', cats='CAUSE.CATEGORY')
    results.append(tvd)

results[:5]
```

C:\Users\phuro\AppData\Local\Temp\ipykernel_3788\3431769426.py:7: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Out[]: [0.02090156380701231,
0.04463513532672648,
0.039967107595026426,
0.030191541547927924,
0.028766676591099202]

```
In [ ]: # Get p-value
```

```
pval = np.mean(np.array(results) >= observed_tvd)
pval
```

```
Out[ ]: 0.0014
```

```
In [ ]: # Plot observed tvd and simulated tvds
```

```
hyp1 = px.histogram(pd.DataFrame({'simulated tvds': results}), \
    x='simulated tvds', histnorm='probability',
    title='Empirical Distribution of TVD')
hyp1.add_vline(x=observed_tvd, line_color='rgb(0,100,80)',
    annotation_text='observed', annotation_position='top right')
p_95 = np.percentile(results, 95)
hyp1.add_vrect(x0=p_95, x1=np.max(results), line_color='rgb(0,176,246)',
    annotation_text='significance level', annotation_position='top left')
hyp1.update_layout(xaxis_range=[0, 0.11])
```

```
In [ ]: # Write to html file
```

```
path = os.path.join('assets', 'hyp1_results.html')
hyp1.write_html(path, include_plotlyjs='cdn')
```

Conclusion: Reject the null hypothesis at a 5% significance level.

The distribution of cause categories in outages with extreme anomaly levels are statistically significantly **different** than outages with regular anomaly levels